

Optimizing LangChain Pipelines for Trust, Observability, and Compliance

Audit-Ready RAG for Banking and Regulated Domains

Sivakumar Mahalingam



Why Trust, Compliance, and Observability Matter?



TRUST

Prevent hallucinations
& enforce citations



COMPLIANCE

Mask PII, follow jurisdiction
rules, avoid policy violations



OBSERVABILITY

Full traceability for
audits and debugging

Optimization Importance



The Banking Challenge

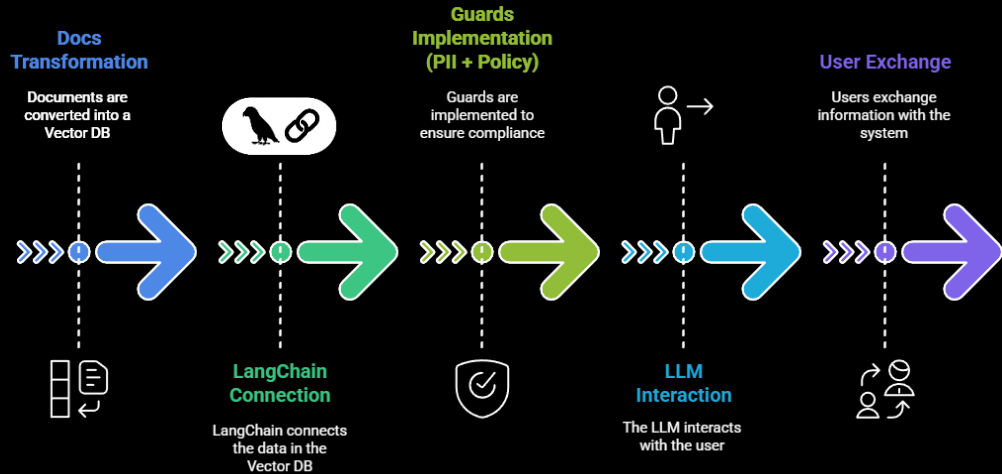
- Analysts and agents need fast answers (KYC, AML, fees)
- Knowledge is scattered: PDFs, policies, SOPs, product docs
- Wrong or non-compliant answers = regulatory risk



Naive pipelines suffer from **latency**, hallucinations, and lack of auditability, impacting reliability and user confidence.

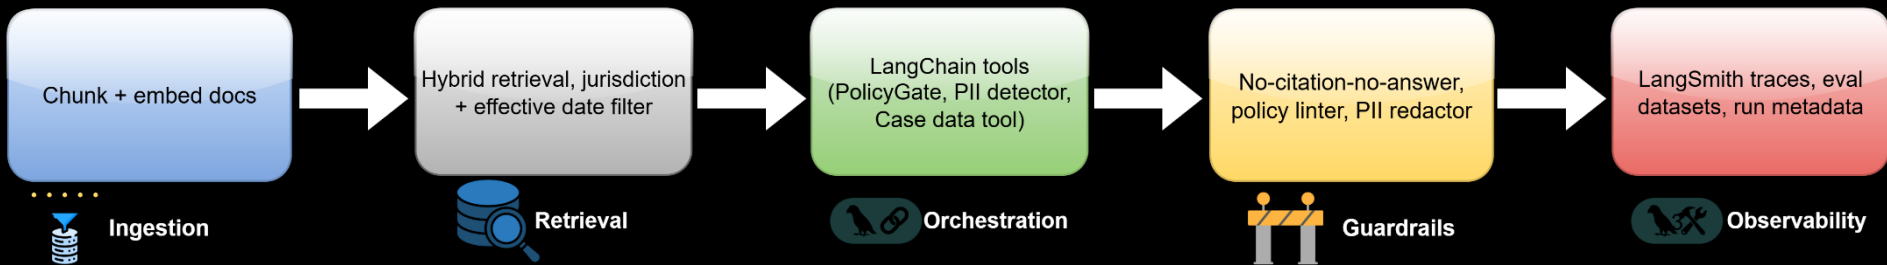
Audit-Ready RAG Copilot

- Retrieval-Augmented Generation with guardrails
- Jurisdiction-aware policy enforcement
- Full LangSmith traceability & replay



This is not just RAG, it's **RAG with governance**
Ensures every answer is backed by citations

Architecture Overview



Code Repository Layout

```
app/  
├── chains/      # LangChain graphs & tools  
├── guards/     # PII, policy linter  
├── retrieval/   # Indexer, retriever  
├── policies/   # Policy packs  
├── schemas/    # Pydantic models  
├── server/     # FastAPI endpoints  
└── telemetry/  # LangSmith client
```




How Guardrails Work?

- **PII Redaction:** Masks PAN, IBAN, SSN, EID pre + post generation.
- **Policy Linter:** Blocks banned phrases like “guaranteed approval”
- **No Citation, No Answer:** Refusal if no supporting docs

Example:

- Input: “4242 4242 4242 1234”
- Output: “**** * 1234”

Guardrails — the backbone of trust and compliance in this system

Observability with LangSmith?

- Every run logged with:
 - Inputs, retrieved docs, tool calls, outputs
 - Metadata: jurisdiction, policy_pack_version, timestamp
- Replay past runs for audits



LangSmith

THANK YOU

Questions & Discussions